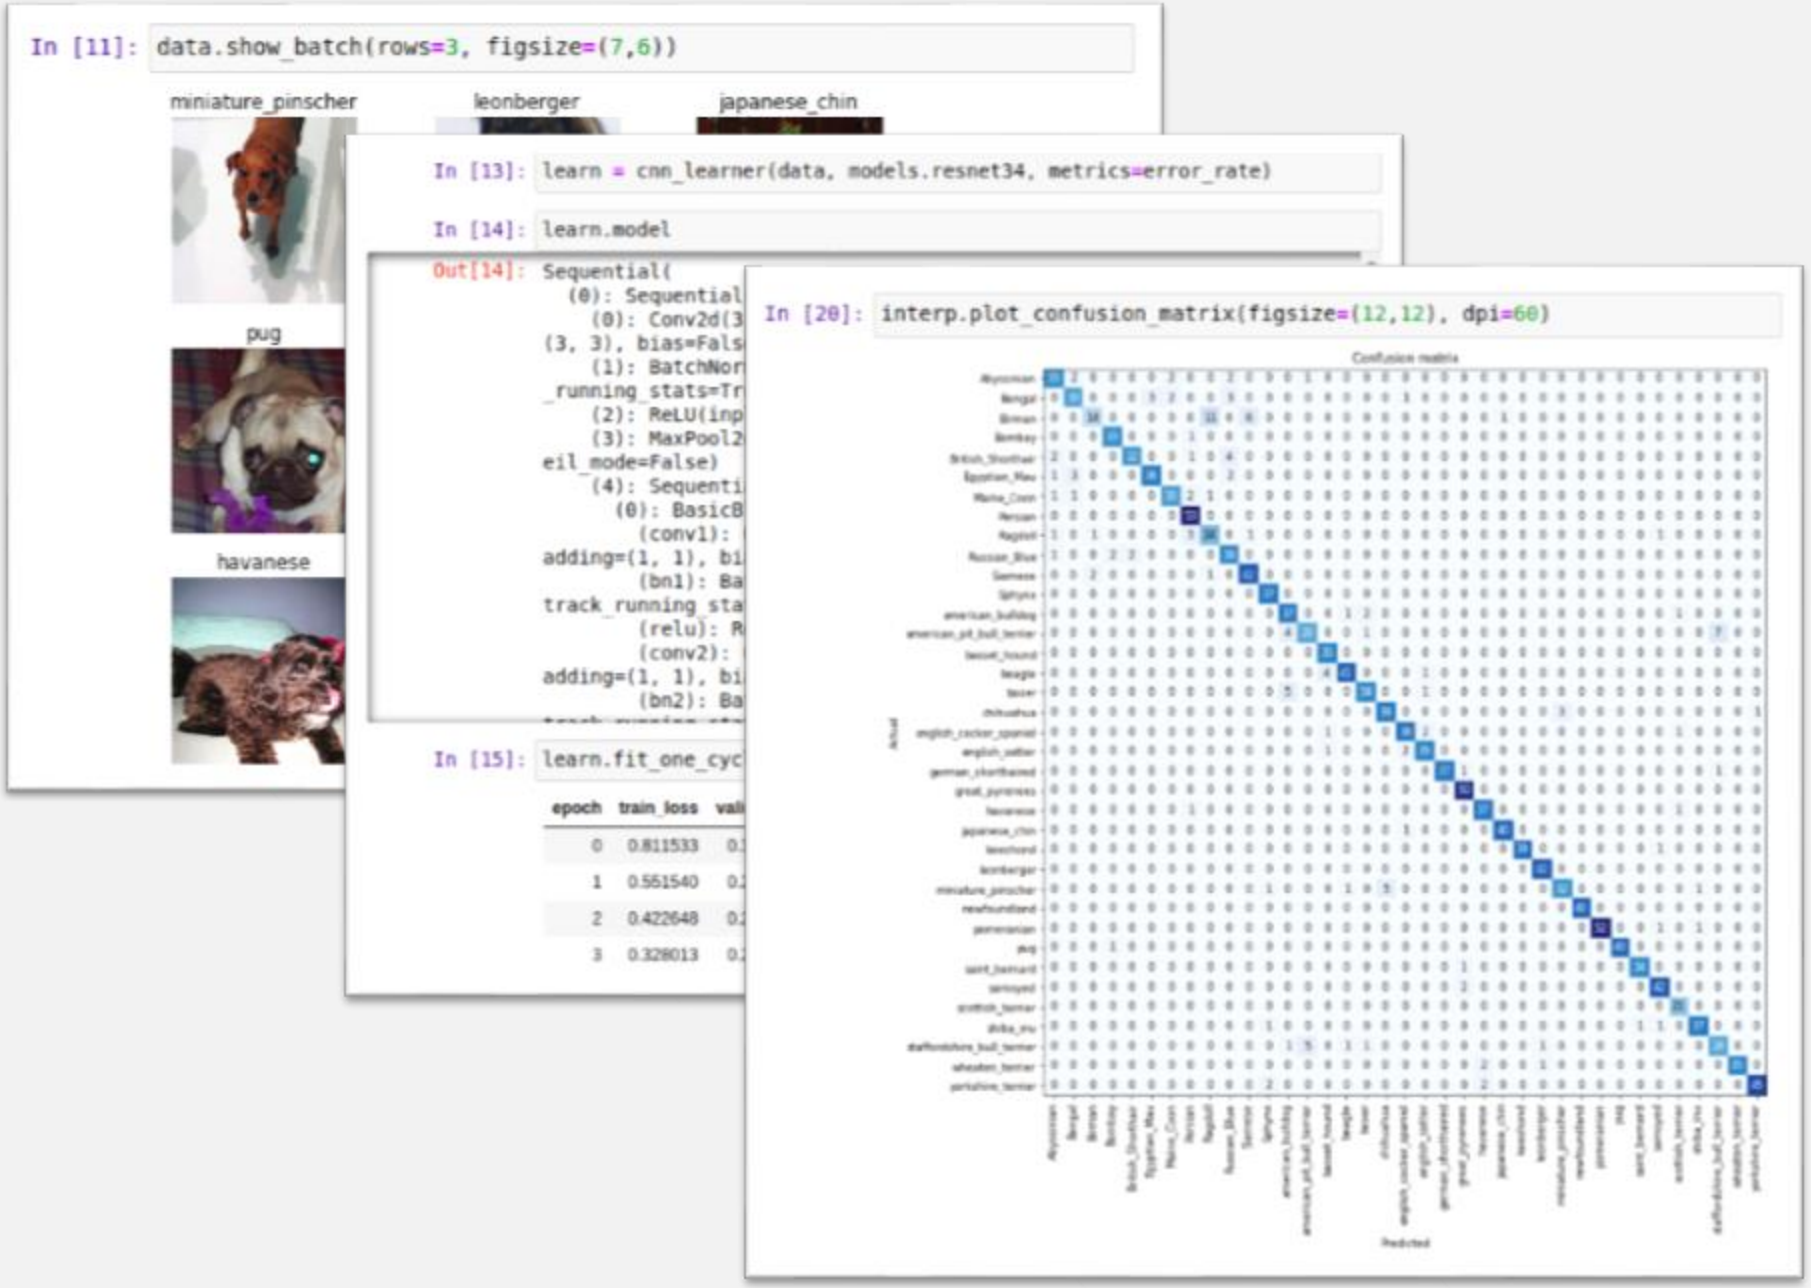

fast.ai

Making neural nets uncool again

CAIR

Josiah Laivins

February 23, 2020

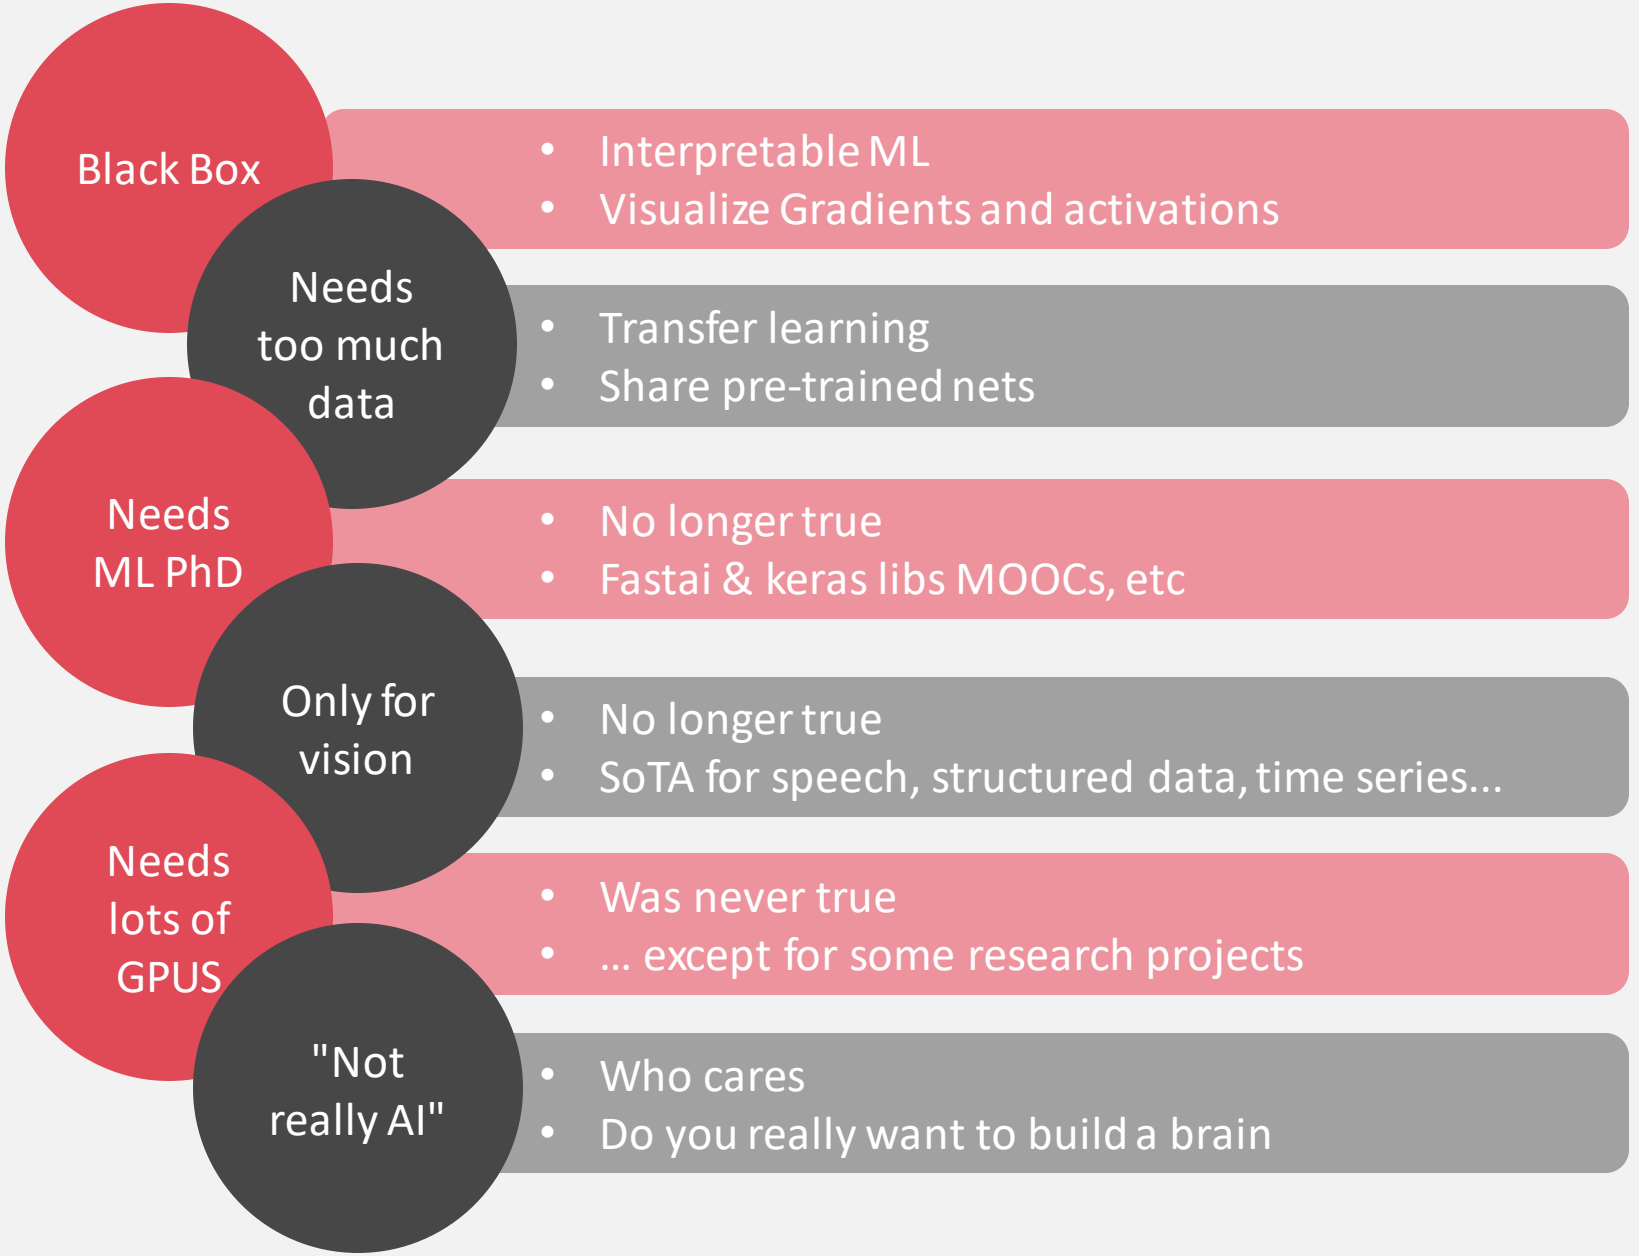

Description from: https://course.fast.ai/videos/?lesson=1

Making Deep Learning Accessible

Software
To make these available to use quickly, reliably, and with minimal code

Community
So that we can all help each other

Research
Ways to make state of the art deep learning techniques more accessible

Research
So that as many people as possible use these

Description from: https://course.fast.ai/videos/?lesson=1

What is fastai?

Popular ML library

One of the most popular programming languages

Different from classical programming. "**friendlier**".

[1] https://pytorch.org/
[2] https://python.org/
[3] https://jupyter.org/

DataBunch

- Images, text, audio
- Data transformation
- Data analysis

Model

- Neural net being trained
- Atomic and portable

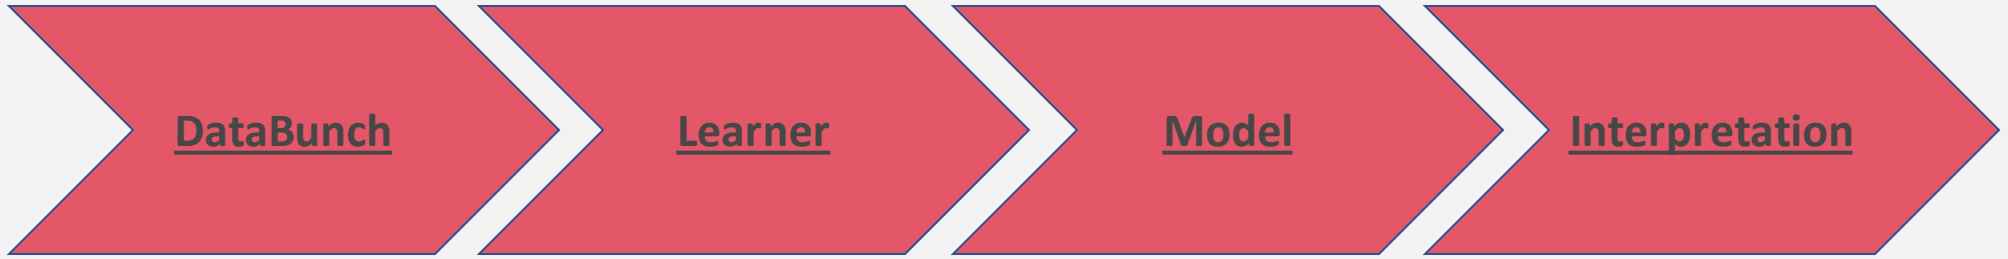

Learner

- Contains everything
- Says how and when things should happen

Interpretation

- Post training quick analysis of model performance, troubleshooting

Who is fastai for?

Newbie's

Tied to a course for anybody to learn from

No math required to get into ML

Jupyter notebooks are far friendlier than most IDEs

Experienced Data Scientists

Existing API code for pipelining

Easy comparison between other models

Attached course goes into detailed topics

Established codebase to build on top of

Description from: https://course.fast.ai/videos/?lesson=1

fastai
vision models

Computer Vision models zoo

The fastai library includes several pretrained models from torchvision, namely:

- resnet18, resnet34, resnet50, resnet101, resnet152
- squeezenet1_0, squeezenet1_1
- densenet121, densenet169, densenet201, densenet161
- vgg16_bn, vgg19_bn
- alexnet

On top of the models offered by torchvision, fastai has implementations for the following models:

- Darknet architecture, which is the base of Yolo v3
- Unet architecture based on a pretrained model. The original unet is described here, the model implementation is detailed in `models.unet`
- Wide resnets architectures, as introduced in this article

[1] https://docs.fast.ai/vision.models.html

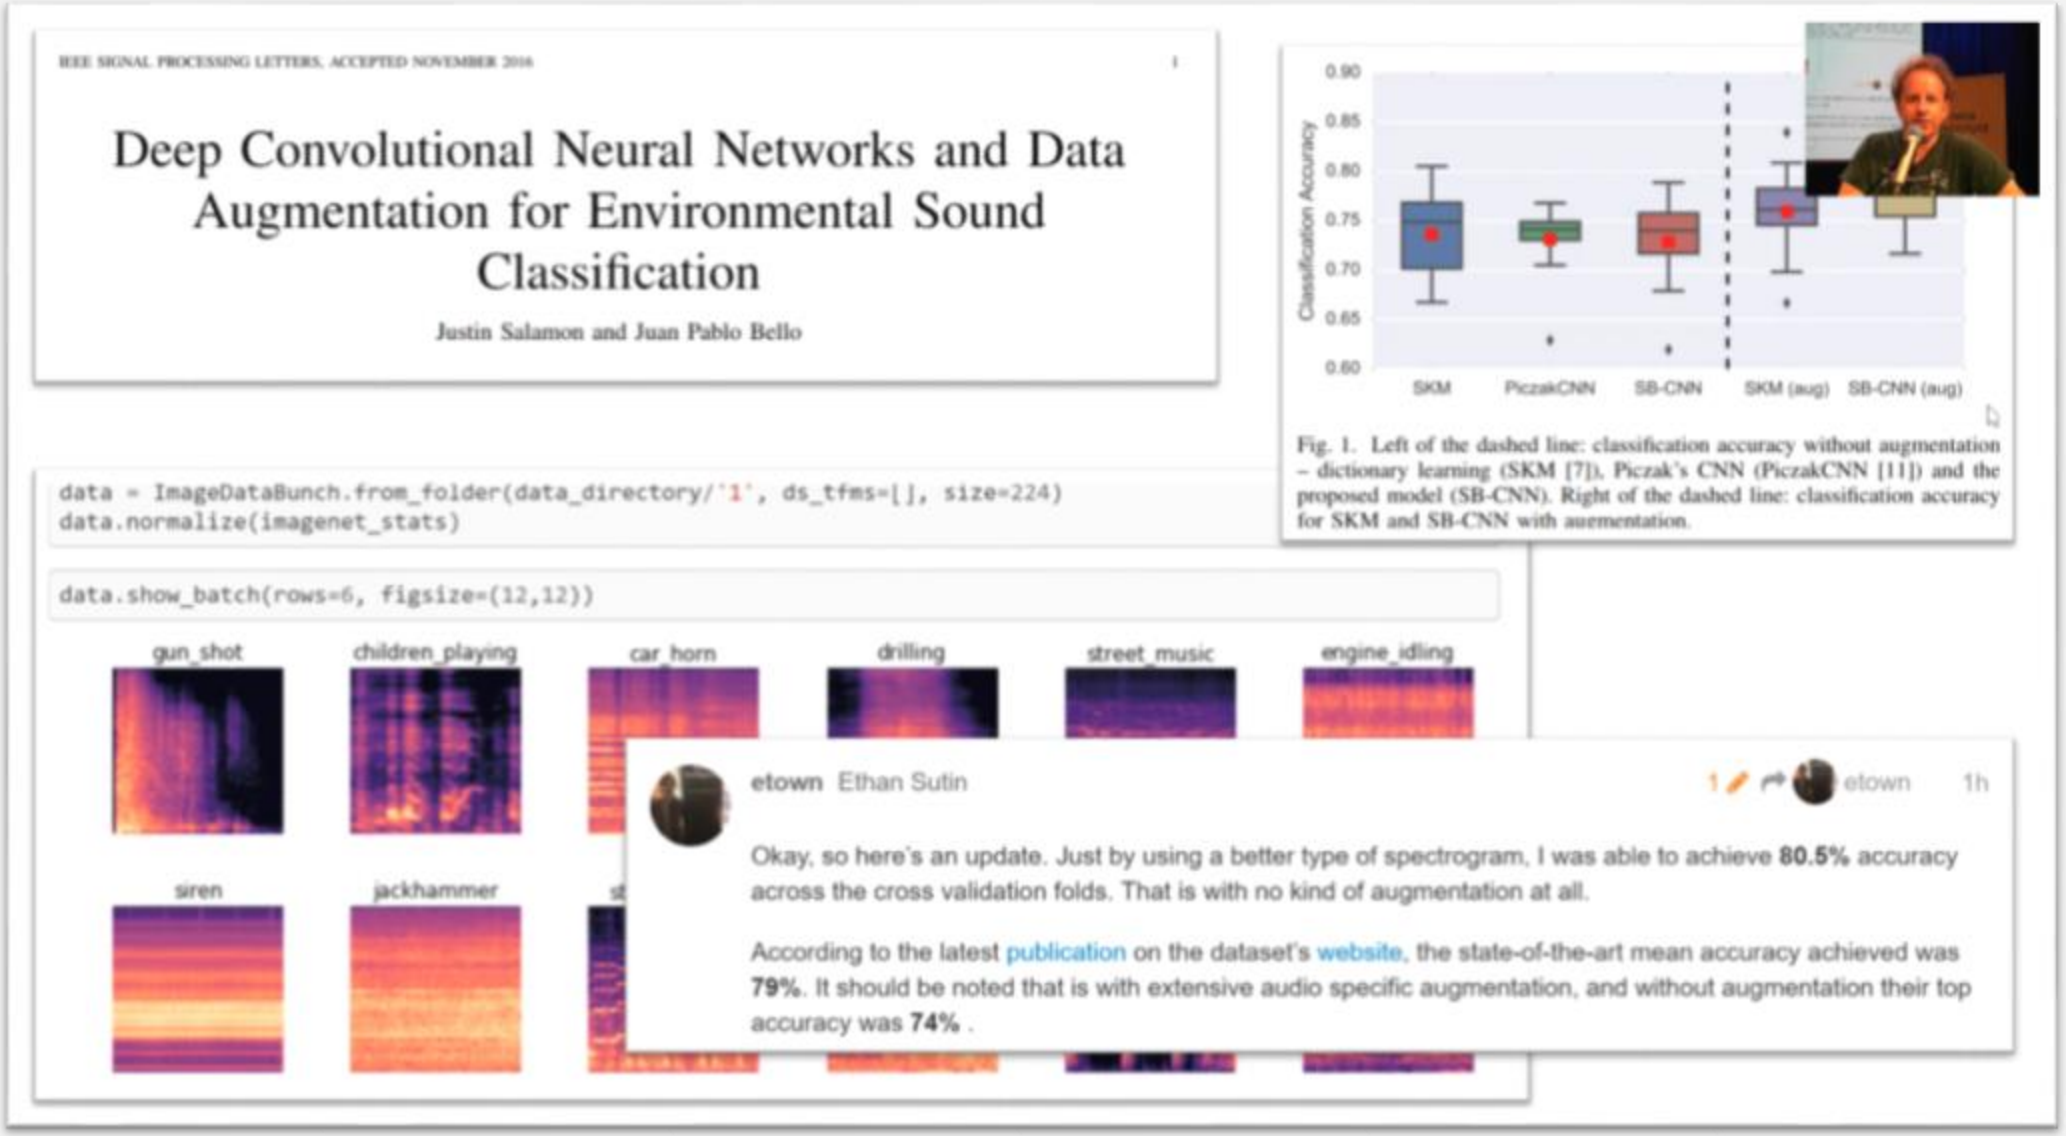

[1]: https://github.com/hiromis/notes/blob/master/Lesson2.md
[2]: https://forums.fast.ai/t/share-your-work-here/27676/215

State of the art on DHCD (देवनागरी)

Suvash Thapaliya @suvash · Oct 27

After lots of intentional tuneups, I can now announce that SoTA for accuracy on #DHCD #देवनागरी dataset is 99.02% (previously 98.50%). All thanks to @fastdotai library & courses, and the dataset creators @pk_gyawali amongst others. #resnet34 #devanagari #deeplearning

lel fitting (transfer learning) on

_loss	valid_loss	accuracy
342	0.037567	0.989493
461	0.035617	0.990072
426	0.034761	0.990290

aining cycle (Accuracy = 99.02

3 12 63

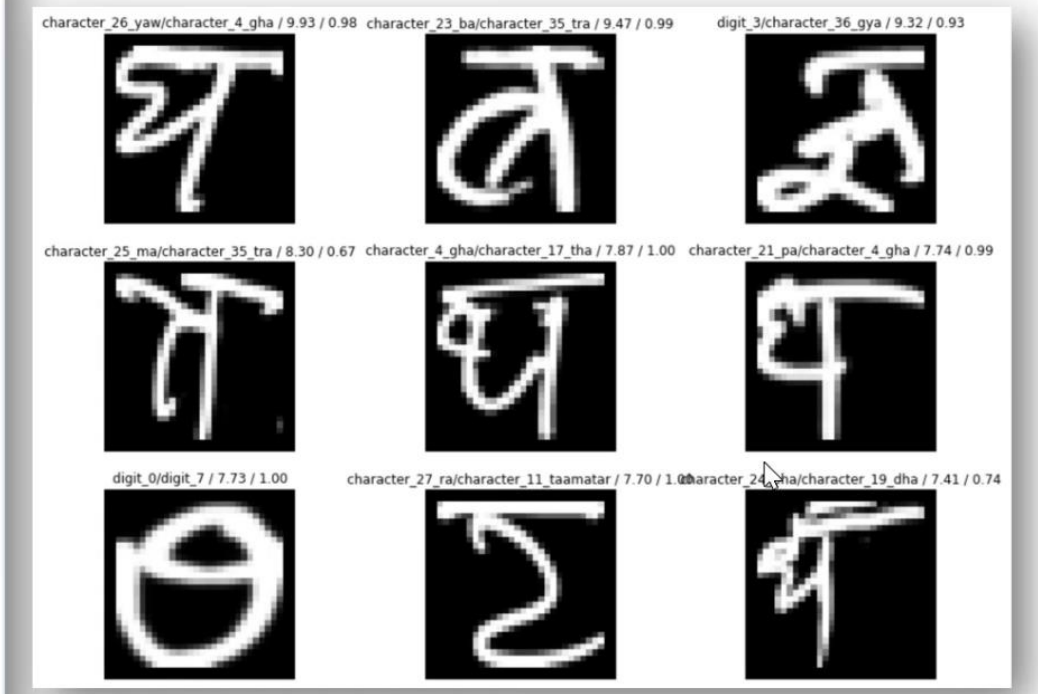

[1]: https://github.com/hiromis/notes/blob/master/Lesson2.md
[2]: https://forums.fast.ai/t/share-your-work-here/27676/38

Guitar Classifier

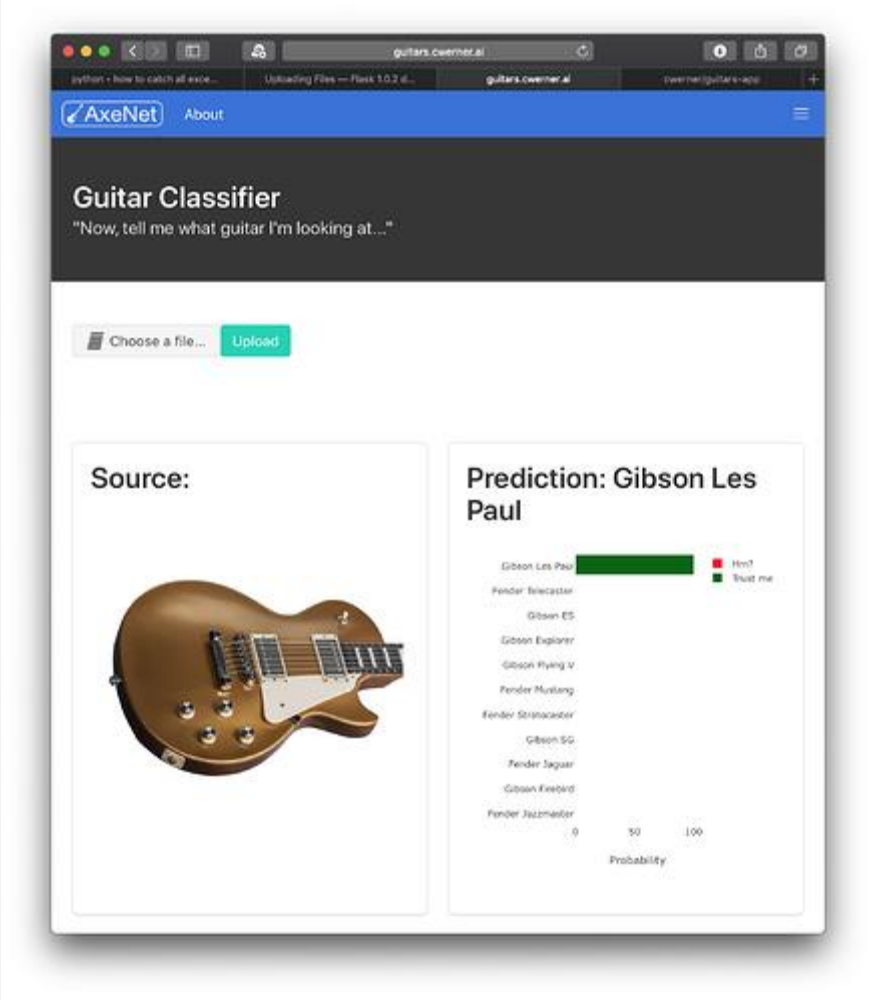

Junk or Health Food Classifier

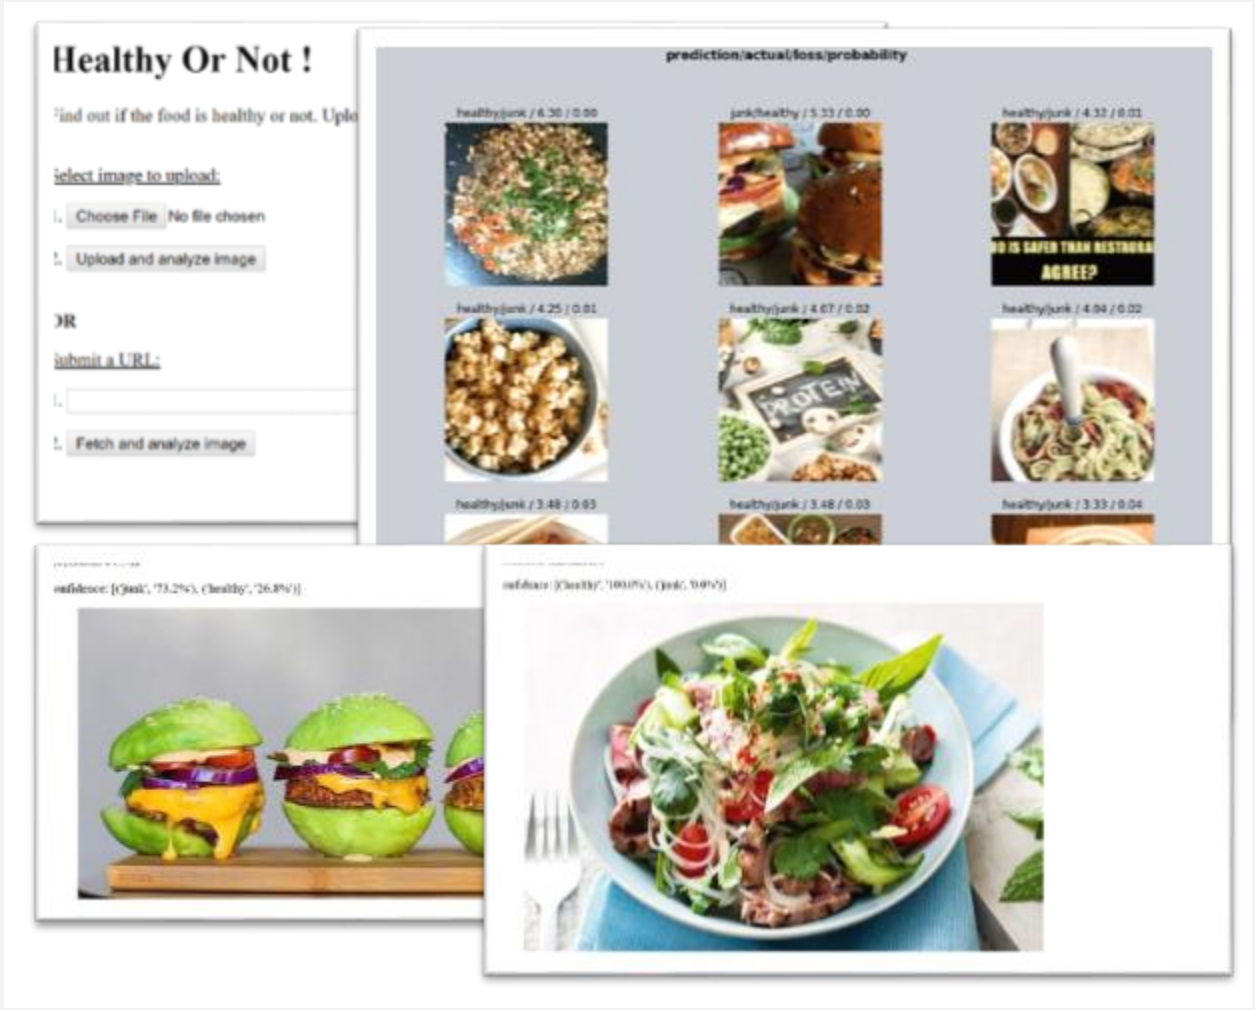

[1]: https://github.com/hiromis/notes/blob/master/Lesson3.md
[2]: https://forums.fast.ai/t/share-your-work-here/27676/399
[2]: https://forums.fast.ai/t/share-your-work-here/27676/340

fast.ai not just classification

NLP

Movie Review Sentiment Classification

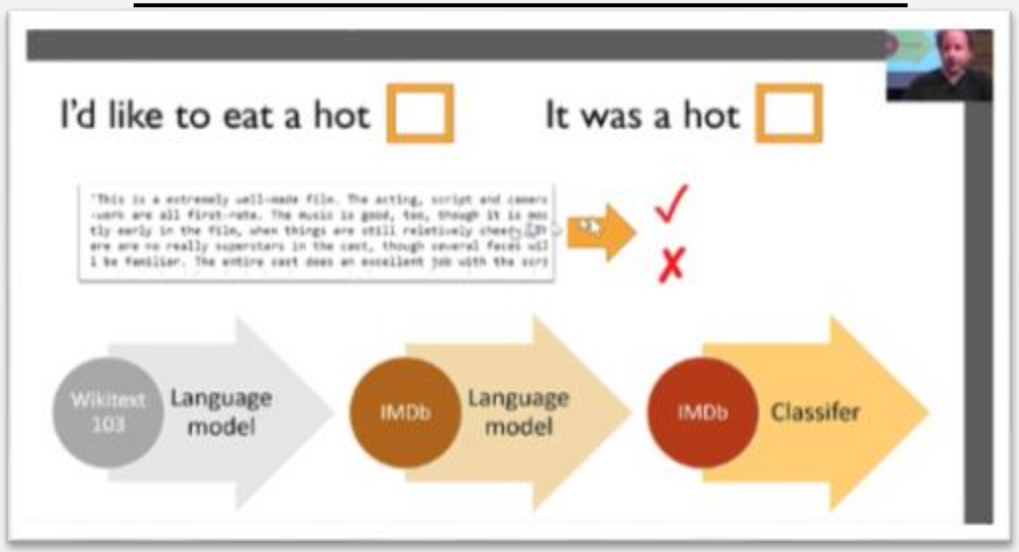

Legal Document Classification

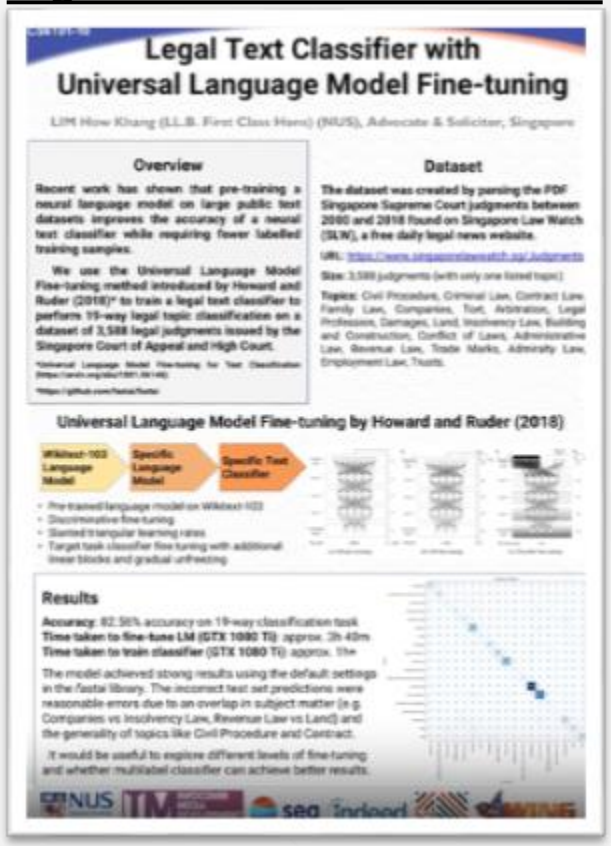

[1]: https://forums.fast.ai/t/deep-learning-lesson-4-notes/30983

GANs

DeOldify

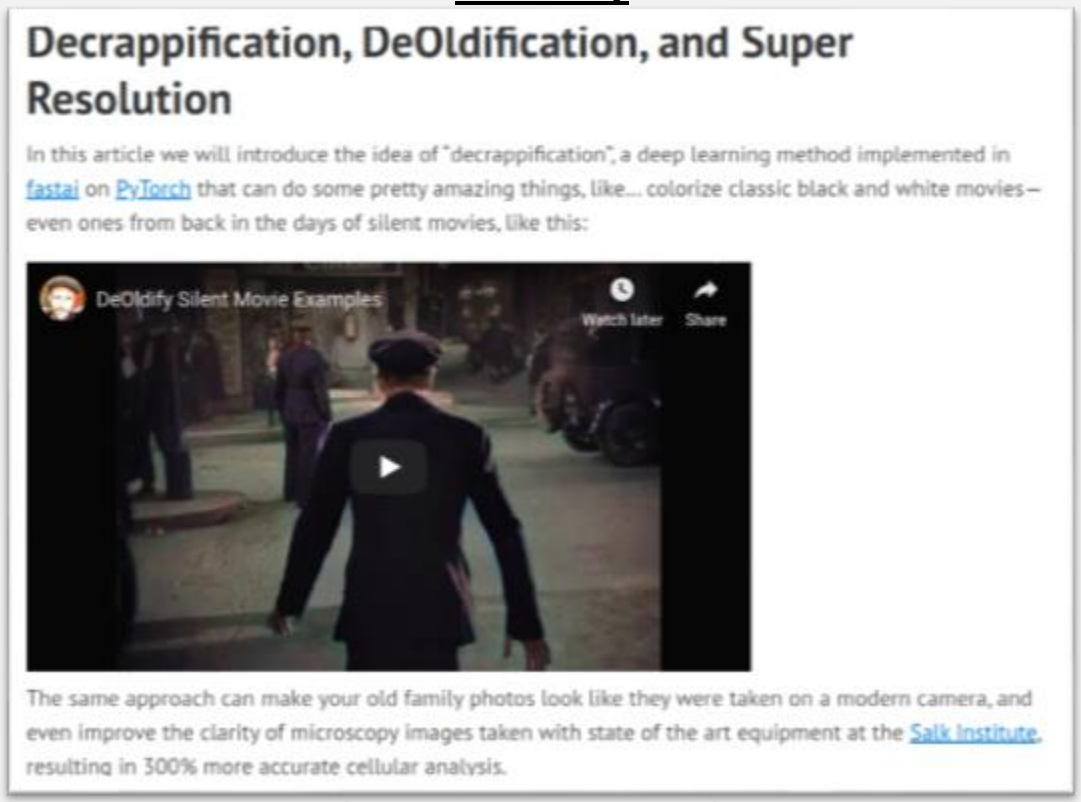

Decrappification, DeOldification, and Super Resolution

In this article we will introduce the idea of "decrappification", a deep learning method implemented in fastai on PyTorch that can do some pretty amazing things, like... colorize classic black and white movies—even ones from back in the days of silent movies, like this:

The same approach can make your old family photos look like they were taken on a modern camera, and even improve the clarity of microscopy images taken with state of the art equipment at the Salk Institute, resulting in 300% more accurate cellular analysis.

Colorizer

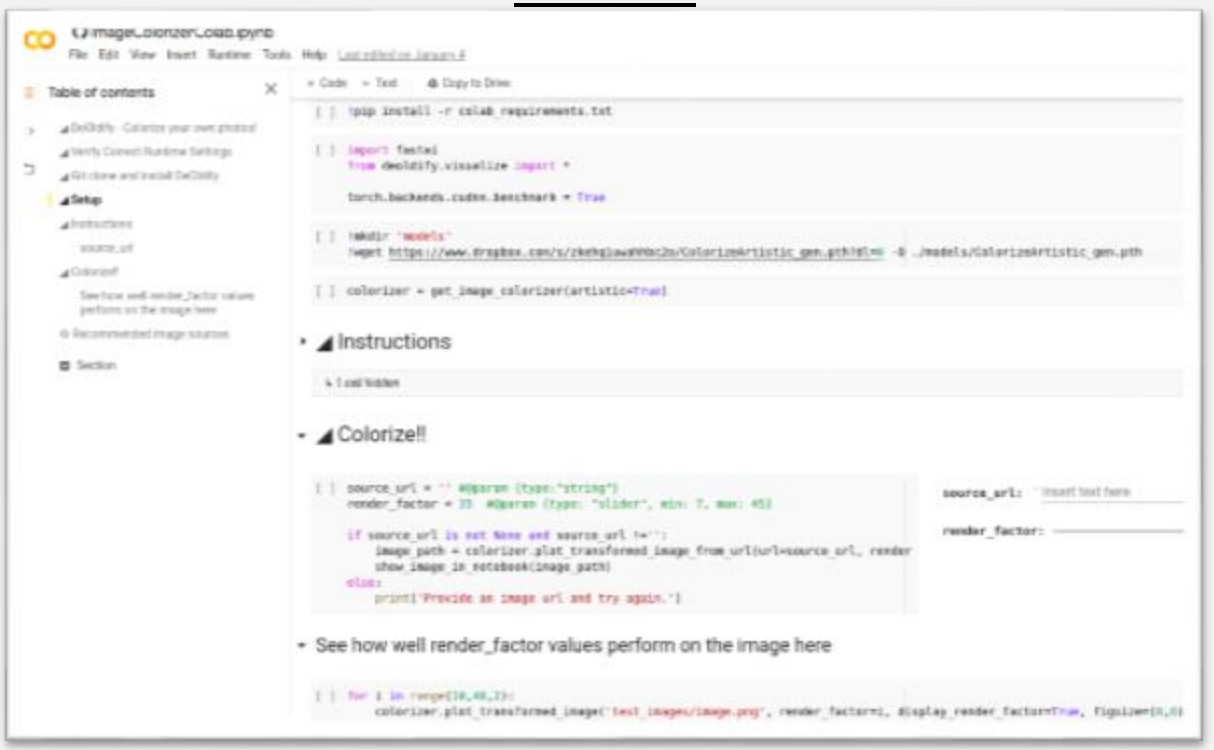

[1]: https://www.fast.ai/2019/05/03/decrappify/
[2]: https://colab.research.google.com/github/jantic/DeOldify/blob/master/ImageColorizerColab.ipynb#scrollTo=LHfUPH42O_iK

Collaborative Filtering

Movie Recommendation System

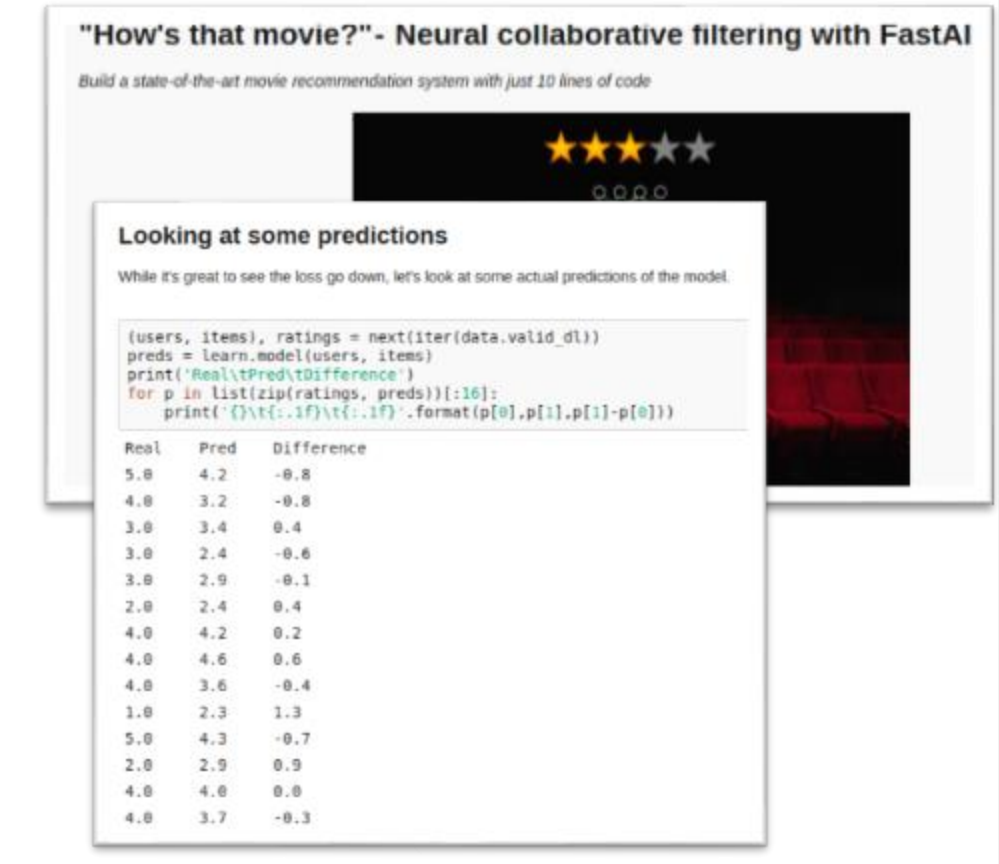

Book Recommendation System

```
1  learn = collab_learner(data, n_factors=40, y_range=(1, 5), wd=1e-1)
```

```
Output:
Top idx:
array(['5000', '3315', '3313', '3312', '3311', '3309', '3308',
'3307', '3306', '3304'], dtype='<U21')

Top names:
array(['Passion Unleashed (Demonica #3)', 'My Story', 'The Gargoyle',
'Pretty Baby', ...,
       'Top Secret Twenty-One (Stephanie Plum, #21)', 'The Warrior
Heir (The Heir Chronicles, #1)', 'Stone Soup',
       'The Sixth Man (Sean King & Michelle Maxwell, #5)'],
dtype='<U144')

Most negative bias:
[(tensor(-0.1021), 'The Almost Moon', 2.49),
 (tensor(-0.0341), 'Skinny Bitch', 2.9),
 (tensor(-0.0325), 'Bergdorf Blondes', 3.0),
 (tensor(-0.0316), 'The Particular Sadness of Lemon Cake', 2.93),
 (tensor(-0.0148), 'The Weird Sisters', 3.08)]
...
```


[1]: https://course.fast.ai/videos/?lesson=4
[2]: https://towardsdatascience.com/collaborative-filtering-with-fastai-3dbdd4ef4f00
[3]: https://jovian.ml/aakashns/5bc23520933b4cc187cfe18e5dd7e2ed

Reinforcement Learning (experimental)

```
In [ ]: memory = ExperienceReplay(memory_size=1000000, reduce_ram=True)
        explore = GreedyEpsilon(epsilon_start=1, epsilon_end=0.1, decay=0.001)
        model = create_dqn_model(data=data, base_arch=DQNModule, lr=0.001, layers=[32,32], opt=optim.RM
        Sprop)
        learn = dqn_learner(data, model, memory=memory, exploration_method=explore,
                            callback_fns=[RewardMetric, EpsilonMetric])
        learn.fit(3)
```

```
In [16]: per_group_interp = GroupAgentInterpretation.from_pickle('data/cartpole_dqn', 'dqn_PriorityExper
         ienceReplay_FEED_TYPE_STATE')
         per_group_interp.add_interpretation(group_interp)
         per_group_interp.plot_reward_bounds(per_episode=True, smooth_groups=10)
```

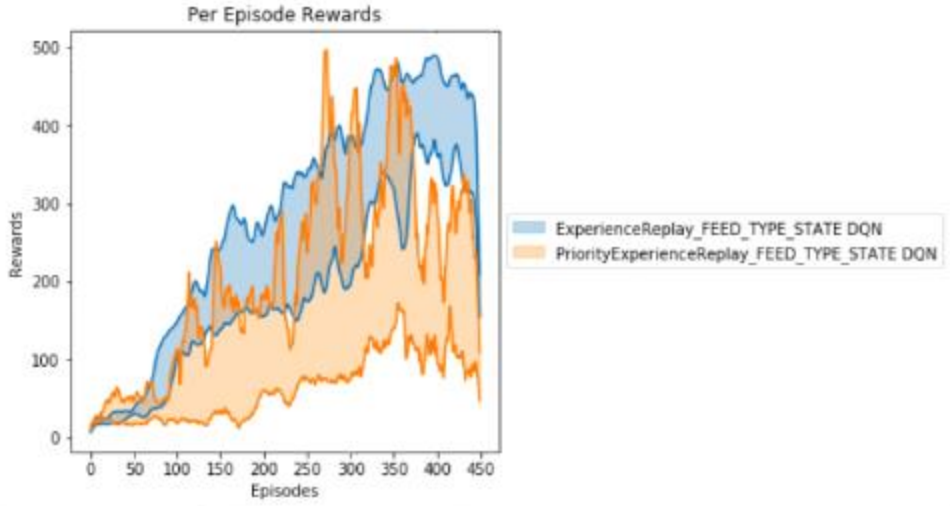

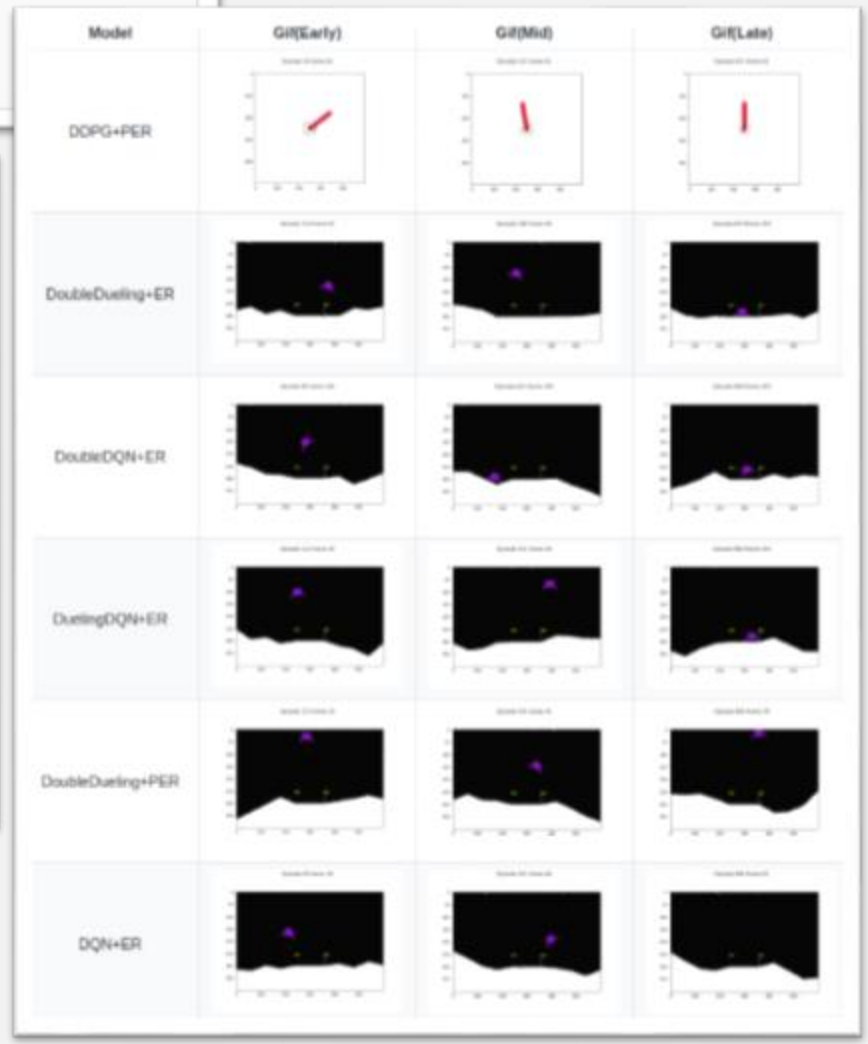

[1]: https://github.com/josiahls/fast-reinforcement-learning

Now to coding

fast.ai has a MOOC!

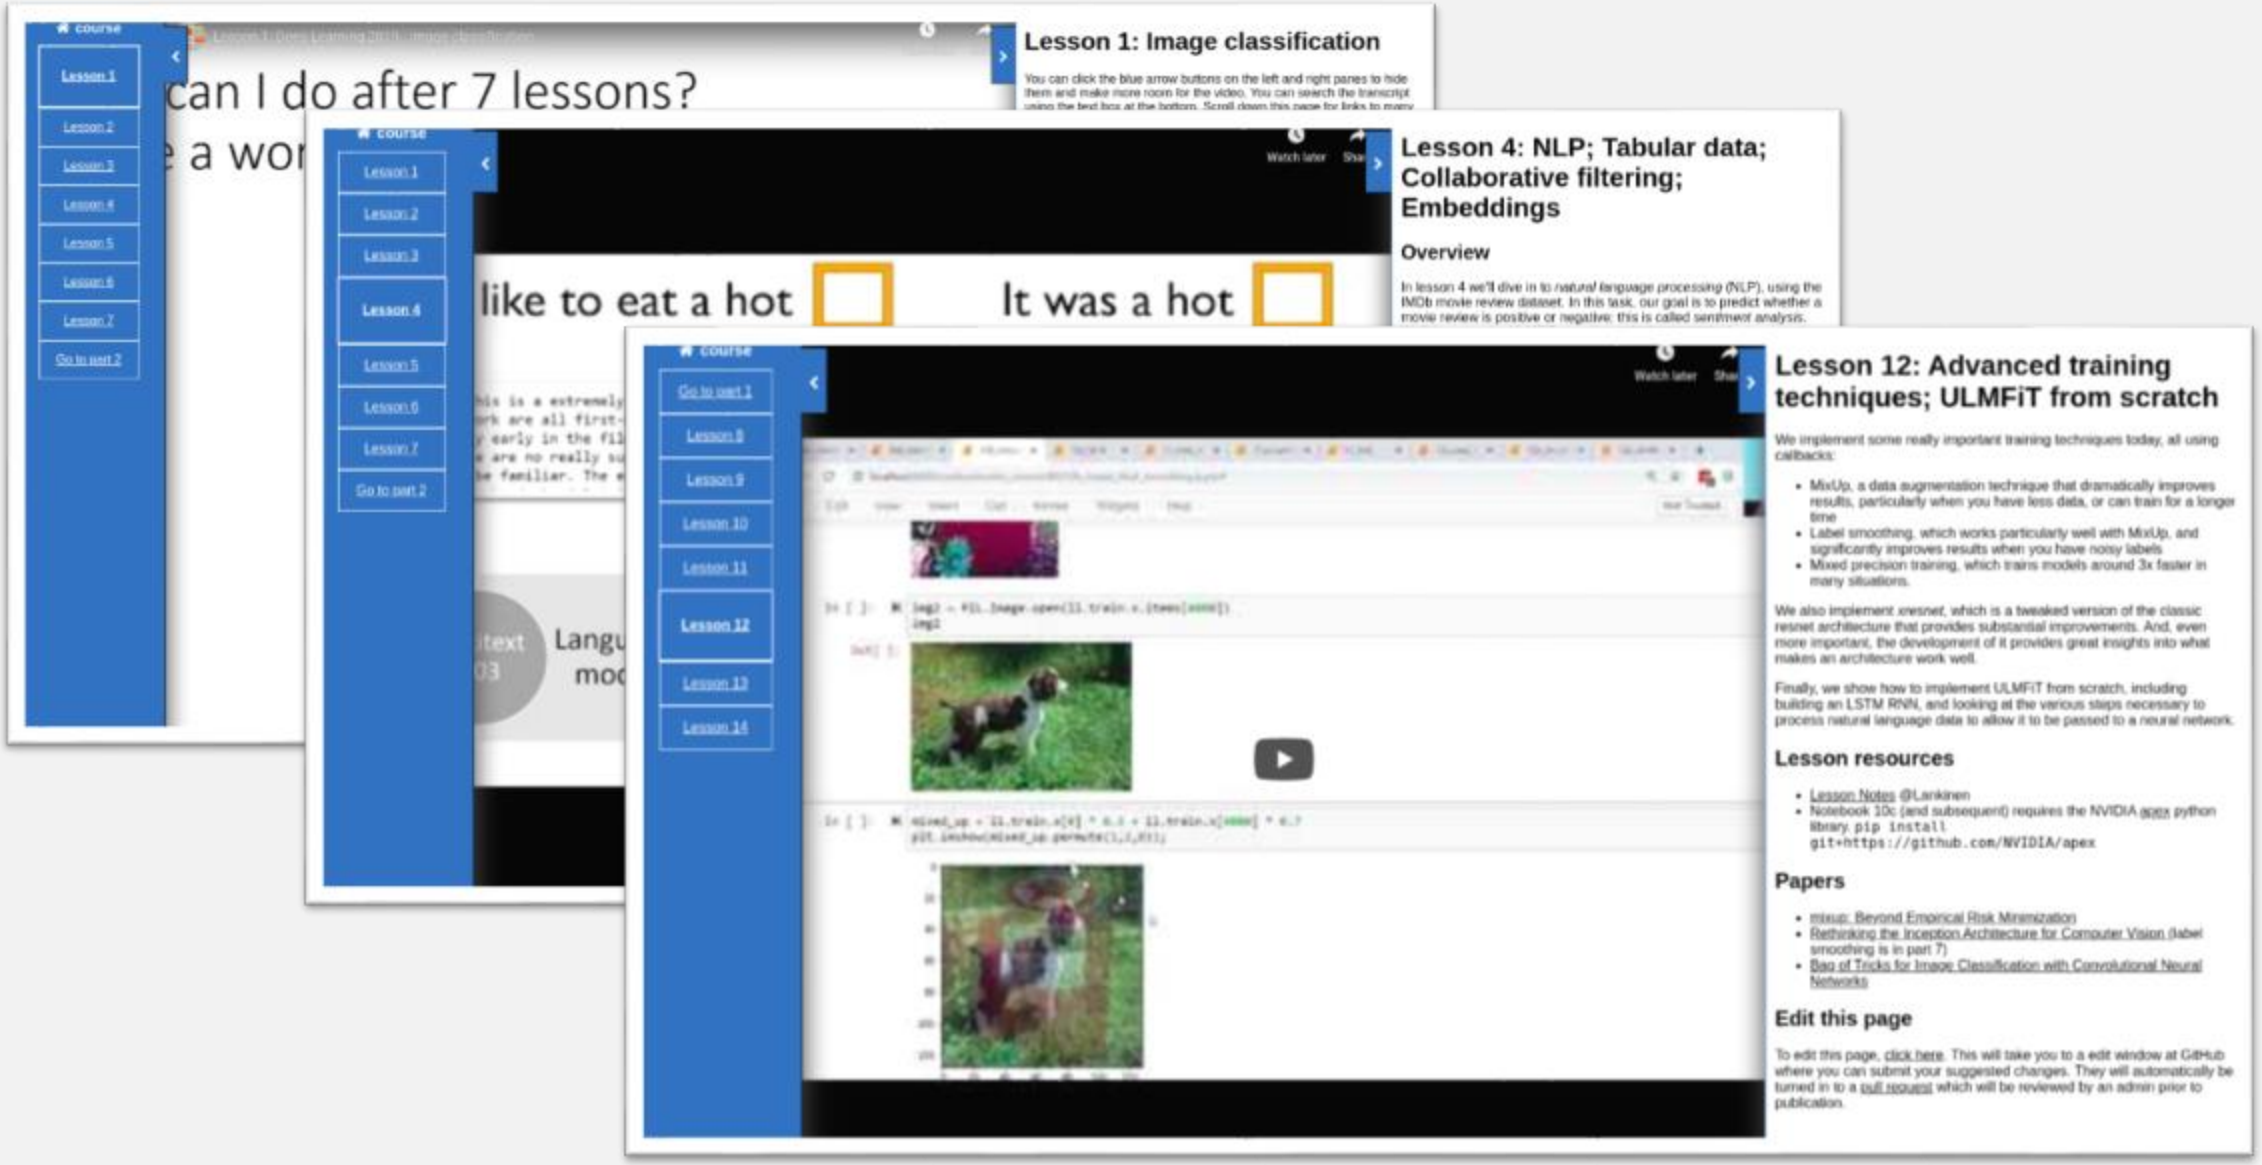

[1]: https://course.fast.ai/videos/?lesson=1

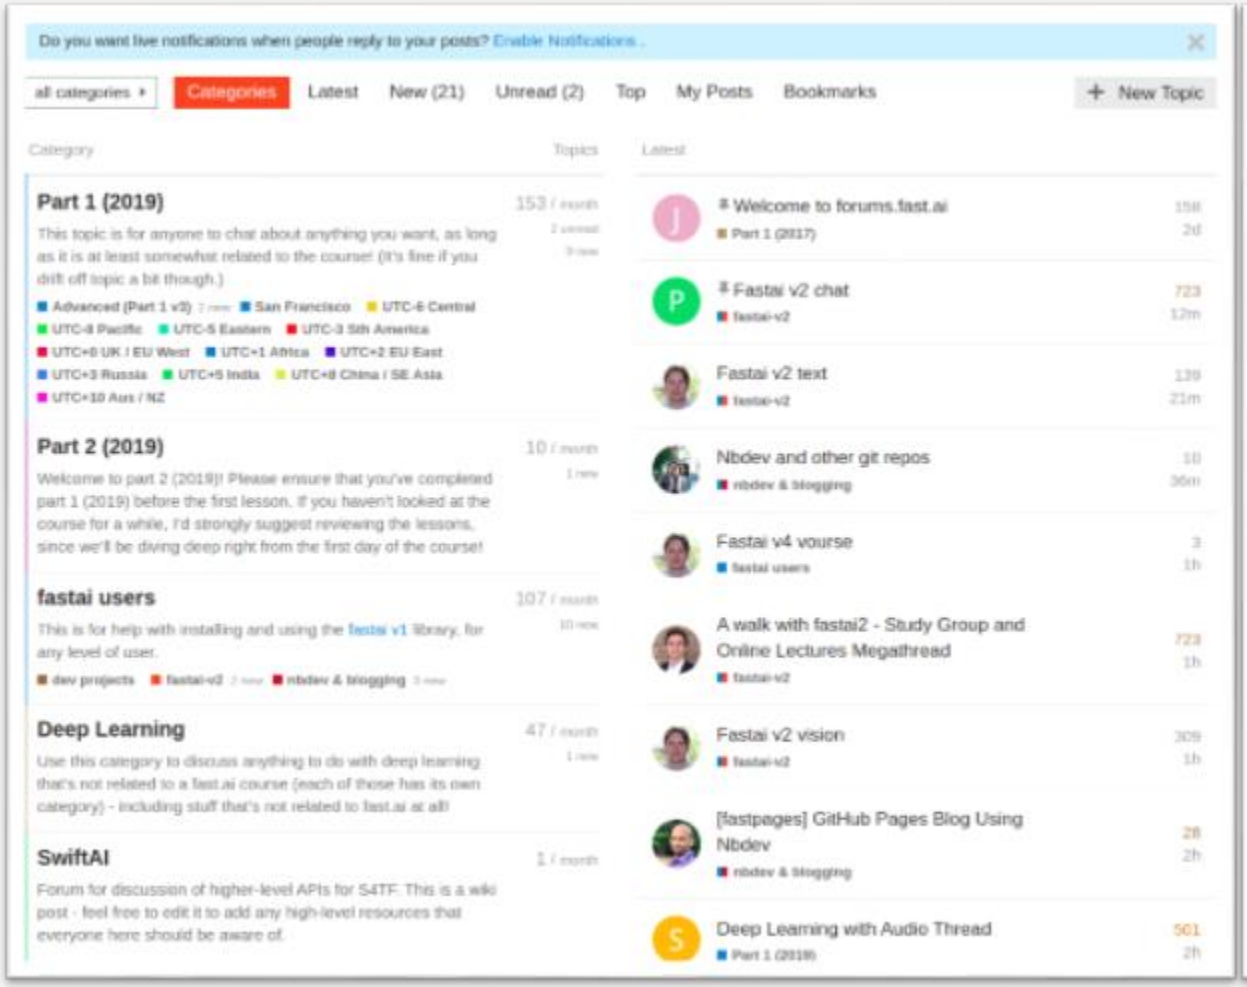

[1]: https://forums.fast.ai/
[2]: https://www.usfca.edu/data-institute/certificates/deep-learning-part-one

*From Stephen Welch's 2/7/2020 CAIR presentation

1. Form a fastai study group, contribute to the fastai library
2. Making kick-ass RL tools, push RL research forward– work with Josiah on fastrl!
3. Unsupervised & one-class learning in manufacturing on the MVTech dataset. [Stephen]
I'm personally offering a $250 bounty on this.
4. Open problem in manufacturing and autonomous driving – supervised model health monitoring
- how do we know when a model is no longer performing well without labeling new data?
5. Create a CAIR Kaggle Team
6. Amazon DeepRacer team – I believe Dr. Shin is forming a team
7. Create high quality educational AI resources (videos, blogs, books)- happy to help as I have time.
8. Summer internship working with Stephen Welch & Josiah at Mariner
9. Summer internship at Atrium working on AI for medical applications

Thank you!